Unit AI-2
Retrieval-Augmented Generation

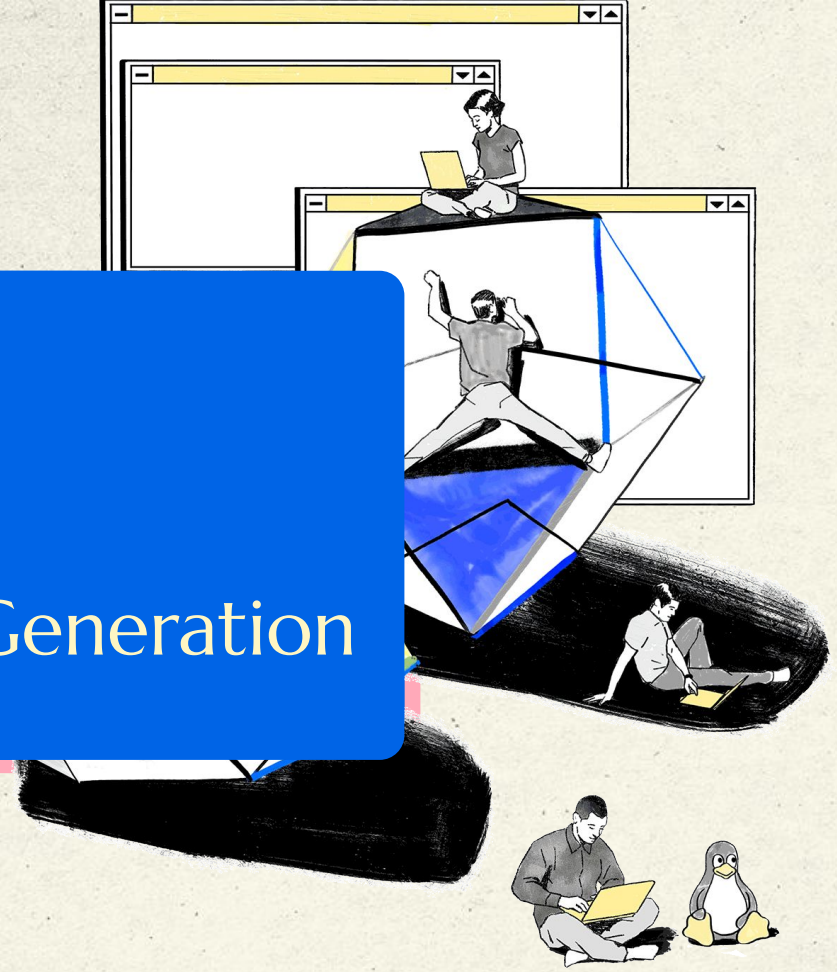

Codesmith

Become Irreplaceable.

Roadmap

- Overview
- How RAG works
- RAG optimizations
- RAG evaluations

Overview

What is RAG?

RAG combines information retrieval with generative AI models to enhance the accuracy and relevance of AI-generated content.

It provides an effective and efficient way of programmatically scaffolding queries with additional context.

Why focus on RAG?

LLMs are not designed to retrieve data - RAG is key to unlocking their potential.

The RAG layer is especially relevant for AI engineering.

Benefits of RAG

- Groundedness
  - Relevance
  - Accuracy
  - Currentness
  - Domain / proprietary specificity
- Interpretability
- Efficiency

Prompting vs RAG vs fine-tuning

Goal	Prompting	RAG	Fine-Tuning
Grounding	😐	✅	😐
Consistency	✅	😐	✅
Confidence	✅	😐	✅
Interpretability	❌	✅	❌
Alignment	😐	😐	✅
Robustness	😐	❌	✅
Latency	❌	😐	✅
Cost	❌	😐	😐

How RAG works

RAG workflow

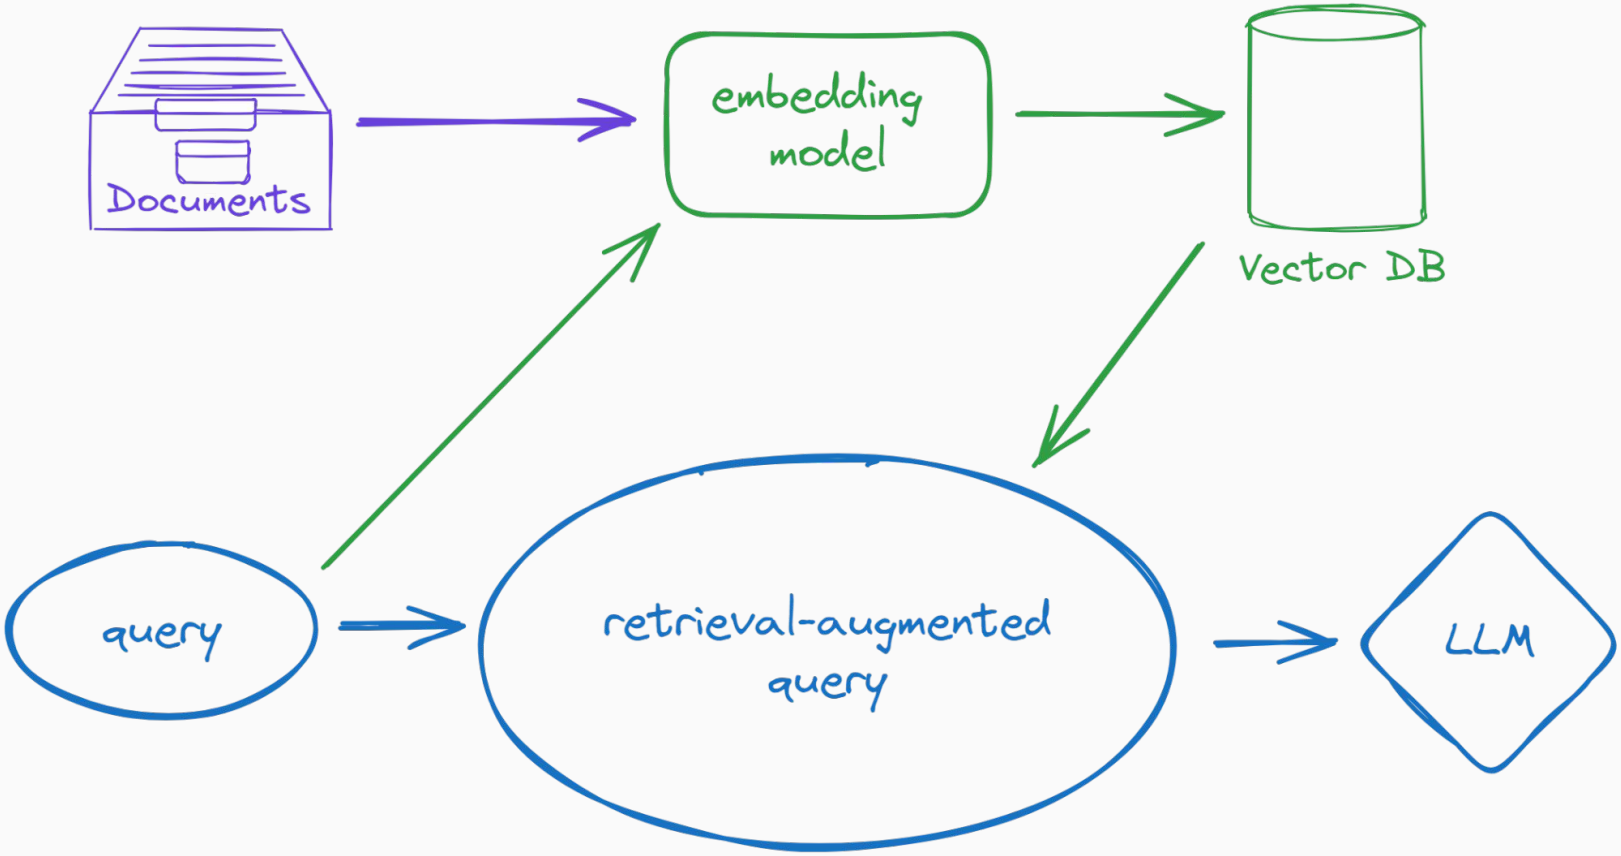

Naive RAG challenges

- Precision
- Recall
- Faithfulness
- Answer relevance
- Accuracy

Data ingestion

- Discovery
- Acquisition
- Validation
- **Transformation**
- **Loading**

RAG vs fine-tuning

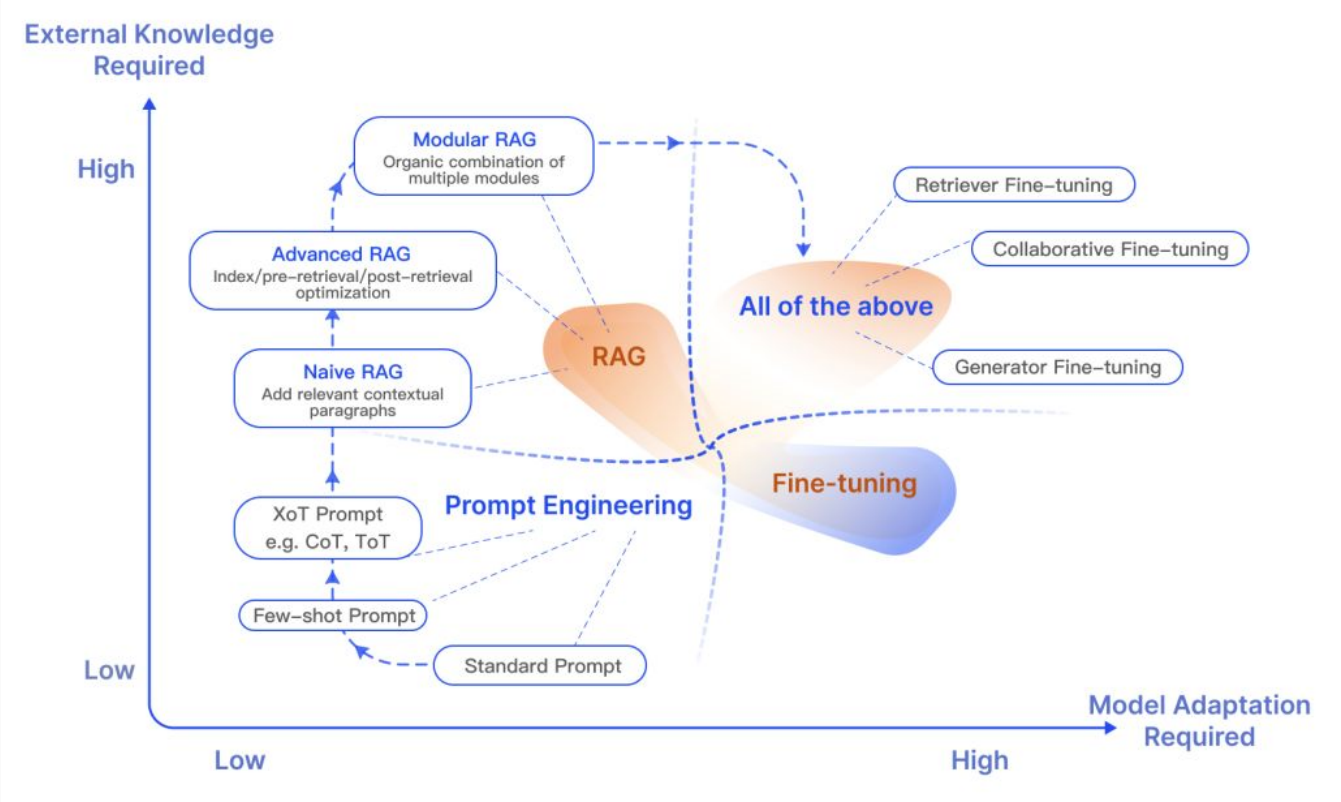

Retrieval-Augmented Generation for Large Language Models: A Survey

Retrieval optimization

Chunk optimization

Balance retrieval specificity against generation context.

- Fixed-size
- Intent-based & Recursive
- Strategy-based

Hybrid retrieval

Capitalize on traditional information retrieval techniques.

- Ensemble (rank aggregation / weighted averaging)
- Cascade (filtering layers)

Keyword + embedding retrieval

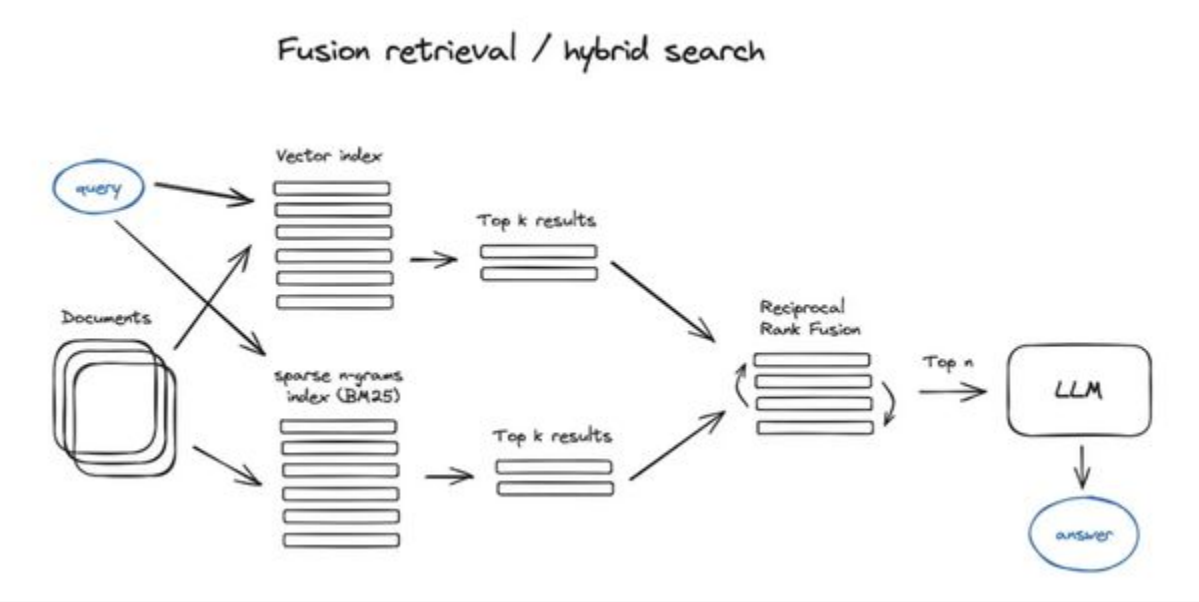

Wei, Huang, and Wang: Retrieval-Augmented Generation for LLM Applications

Metadata filtering

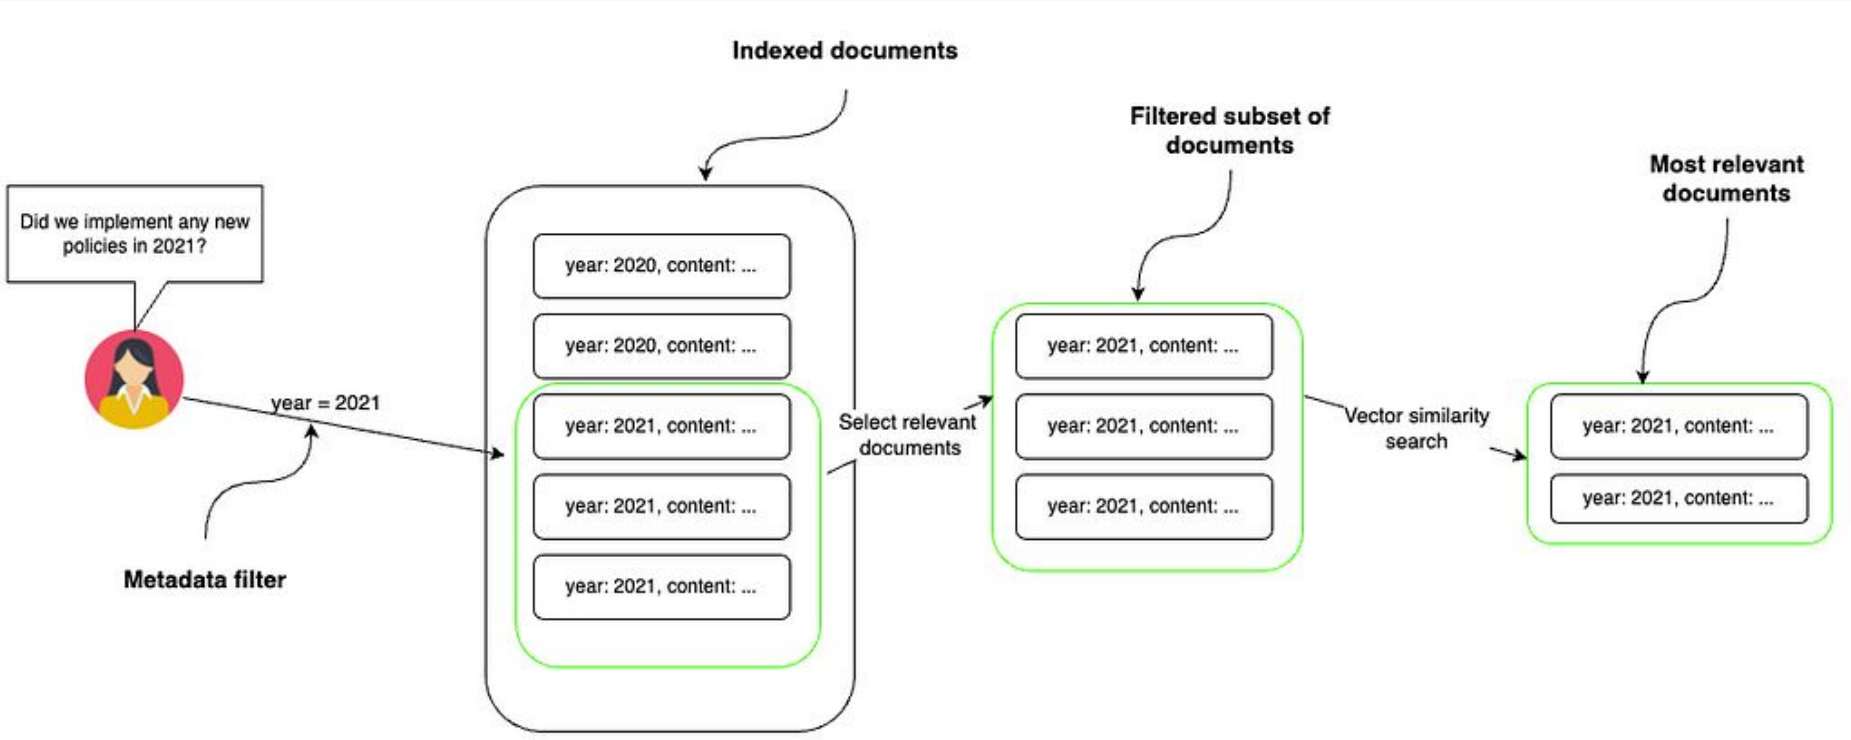

[Neo4j - Graph-based Metadata Filtering to Improve Vector Search in RAG Applications]

Small-to-big retrieval

Optimize for retrieval *and* generation by expanding context after initial search.

- Sentence window retrieval
- Parent document crawler

Similarly, retrieval could proceed recursively to build context.

Parent document crawler

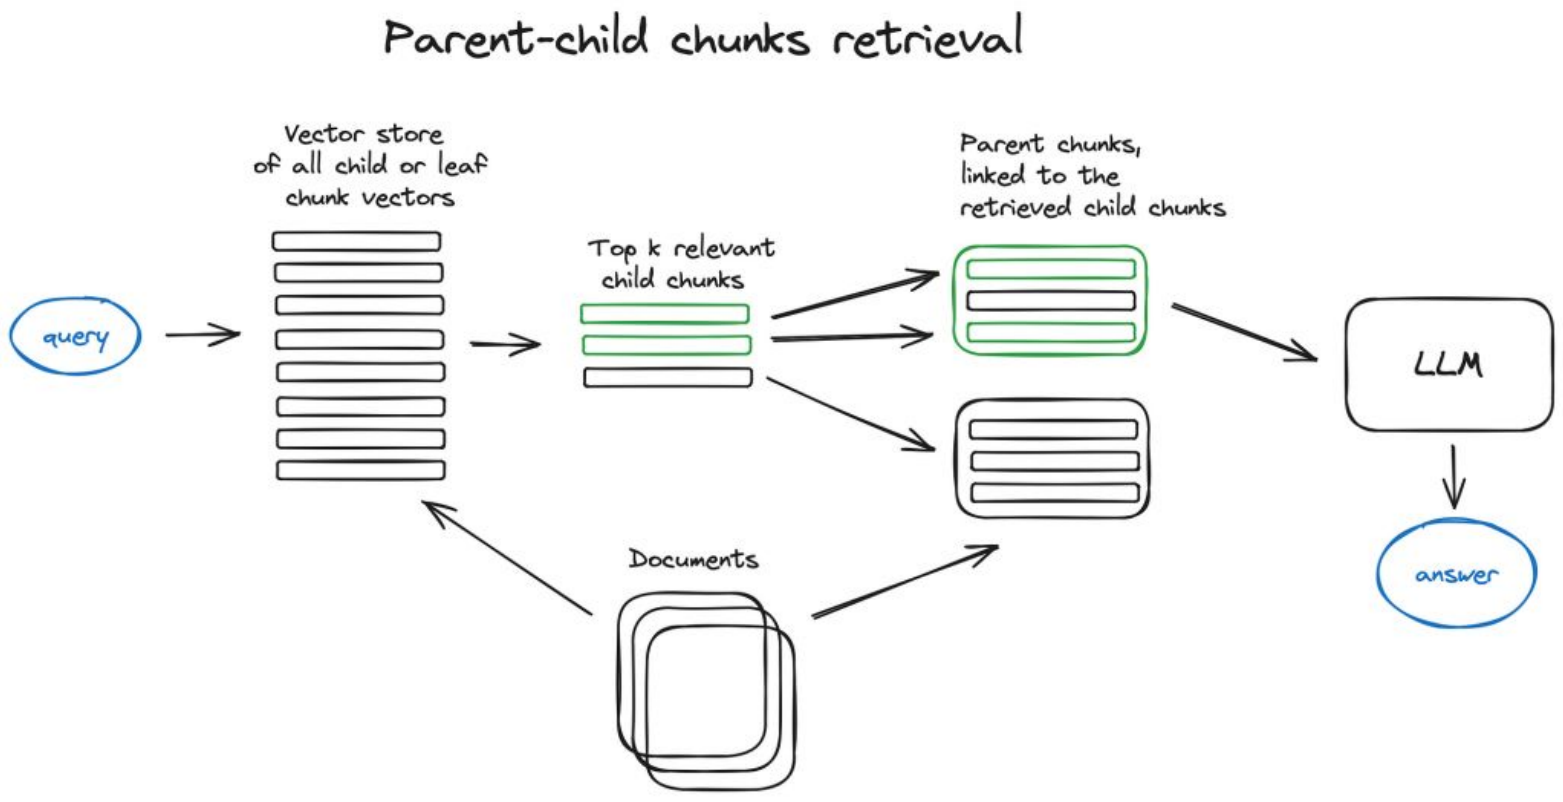

[ChatGen.ai - The Ultimate Guide on Retrieval Strategies - RAG]

Hypothetical Document Embeddings (HyDE)

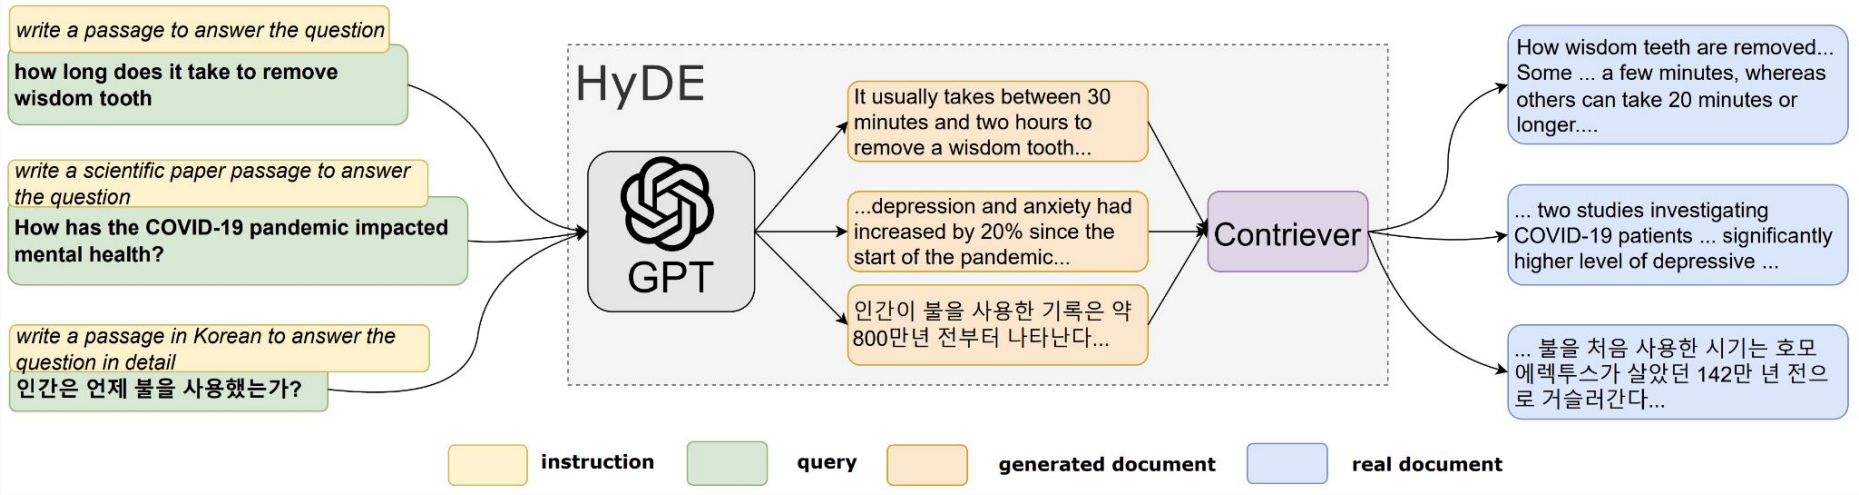

[Precise Zero-Shot Dense Retrieval without Relevance Labels]

And many more...

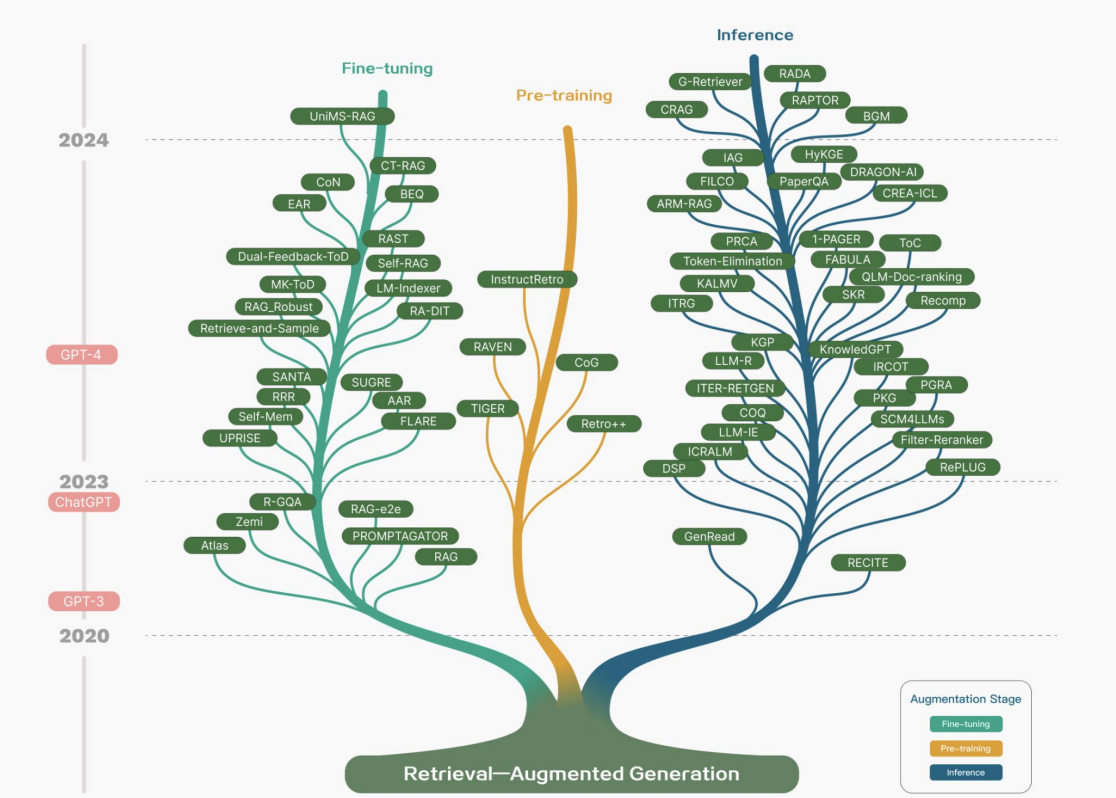

[Retrieval-Augmented Generation for Large Language Models: A Survey]

OpenAI “RAG Success Story”

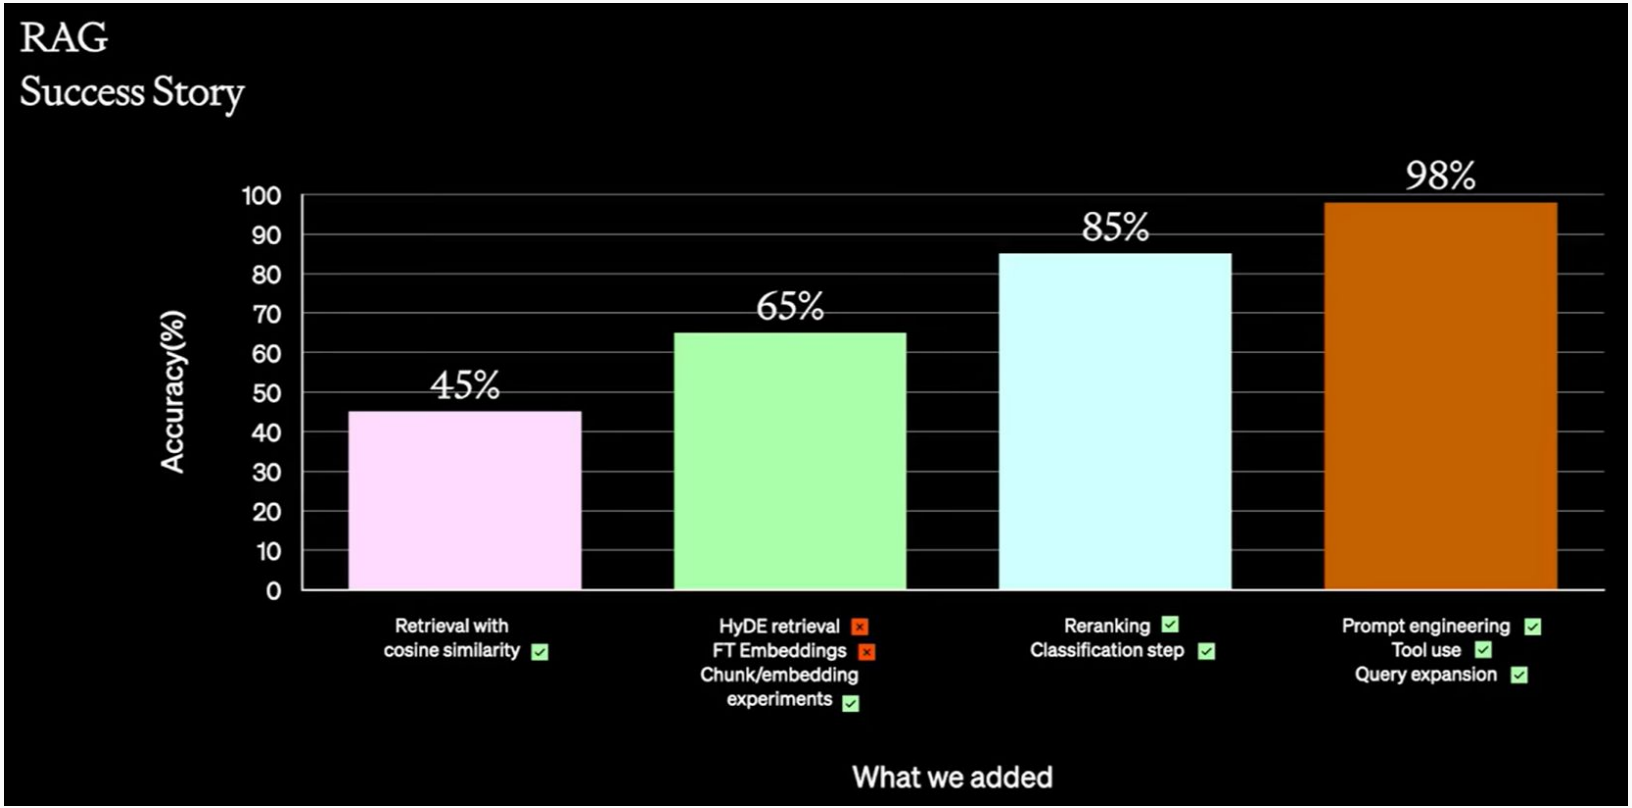

A Survey of Techniques for Maximizing LLM Performance

RAG evaluation

Information retrieval metrics

- Recall
- Precision
- Mean reciprocal rank (MRR)
- Normalized discounted cumulative gain (NDCG)

Building an evaluation dataset

- Benchmarks
- Web data
- Synthetic data
- Proprietary data

RAG testing pyramid

As with traditional applications, it’s beneficial to test each unit (retrieval, augmentation, and generation), their integration, and the end-to-end result.

RAG evaluation dimensions

- Retrieval relevance
- Groundedness
- Adaptability
- Toxicity
- Efficiency

Retrieval relevance

- ROUGE (Recall-Oriented Understudy for Gisting Evaluation)
- RAGAs (RAG Assessment)

Groundedness

- Entity linking accuracy
- Semantic similarity scores
- FactScore

Generation-specific dimensions

- Noise robustness
- Negative rejection
- Information integration
- Counterfactual robustness

Further considerations

RAG vs long context

SOTA models boast long context windows and have dramatically improved their performance - but RAG may still be preferable given:

- Expanded scope
- Increased precision
- Reduced cost

RAGOps

- Vector DBs
- Frameworks + libraries
- RAG 2.0